


## Asynchronous Replication and Bayou

DUKE  
Systems & Architecture

## Asynchronous Replication

**Idea:** build available/scalable information services with *read-any-write-any* replication and a weak consistency model.

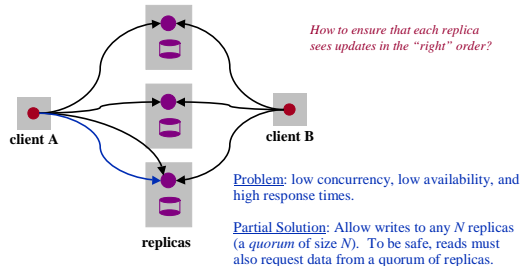
- no denial of service during transient network partitions
- supports massive replication without massive overhead
- "ideal for the Internet and mobile computing" [Golding92]



DUKE  
Systems & Architecture

## Synchronous Replication

**Basic scheme:** connect each client (or front-end) with every replica; writes go to all replicas, but client can read from any replica (*read-one-write-all* replication).



DUKE  
Systems & Architecture

## Grapevine and Clearinghouse (Xerox)

Weakly consistent replication was used in earlier work at Xerox PARC:

- Grapevine and Clearinghouse name services
  - Updates were propagated by unreliable multicast ("direct mail").
- Periodic *anti-entropy* exchanges among replicas ensure that they eventually converge, even if updates are lost.
  - Arbitrary pairs of replicas periodically establish contact and resolve all differences between their databases.
  - Various mechanisms (e.g., MD5 digests and update logs) reduce the volume of data exchanged in the common case.
  - Deletions handled as a special case via "death certificates" recording the delete operation as an update.

DUKE  
Systems & Architecture

## Epidemic Algorithms

PARC developed a family of weak update protocols based on a disease metaphor (*epidemic algorithms* [Demers et. al. OSR 1/88]):

- Each replica periodically "touches" a selected "susceptible" peer site and "infects" it with updates.
  - Transfer every update known to the carrier but not the victim.
  - Partner selection is randomized using a variety of heuristics.
- Theory shows that the epidemic eventually infects the entire population with high probability (assuming it is connected).
  - Probability that replicas that have not yet converged decreases exponentially with time.
  - Heuristics (e.g., push vs. pull) affect traffic load and the expected time-to-convergence.

DUKE  
Systems & Architecture

## How to Ensure That Replicas Converge

1. Using any form of epidemic (randomized) anti-entropy, all updates will (eventually) be known to all replicas.
2. Imposing a global order on updates guarantees that all sites (eventually) apply the same updates in the same order.
3. Assuming conflict *detection* is deterministic, all sites will detect the same conflicts.
  - Write conflicts cannot (generally) be detected when a site *accepts* a write; they appear when updates are *applied*.
3. Assuming conflict *resolution* is deterministic, all sites will resolve all conflicts in exactly the same way.

DUKE  
Systems & Architecture

## Issues and Techniques for Weak Replication

1. How should replicas choose partners for anti-entropy exchanges?  
Topology-aware choices minimize bandwidth demand by “flooding”, but randomized choices survive transient link failures.
2. How to impose a global ordering on updates?  
logical clocks and delayed delivery (or delayed commitment) of updates
3. How to integrate new updates with existing database state?  
Propagate updates rather than state, but how to detect and reconcile conflicting updates? Bayou: user-defined checks and merge rules.
4. How to determine which updates to propagate to a peer on each anti-entropy exchange?  
vector clocks or vector timestamps
5. When can a site safely commit or stabilize received updates?  
receiver acknowledgement by vector clocks (TSAE protocol)

DUKE  
Systems & Architecture

## Bayou Basics

1. Highly available, weak replication for mobile clients.  
Beware: every device is a “server”... let's call 'em sites.
2. Update conflicts are detected/resolved by rules specified by the application and transmitted with the update.  
interpreted dependency checks and merge procedures
3. Stale or tentative data may be observed by the client, but may mutate later.  
The client is aware that some updates have not yet been confirmed.  
“An inconsistent database is marginally less useful than a consistent one.”

DUKE  
Systems & Architecture

## Clocks

1. physical clocks  
Protocols to control drift exist, but physical clock timestamps cannot assign an ordering to “nearly concurrent” events.
2. logical clocks  
Simple timestamps guaranteed to respect causality: “A's current time is later than the timestamp of any event A knows about, no matter where it happened or who told A about it.”
3. vector clocks  
Order(N) timestamps that say exactly what A knows about events on B, even if A heard it from C.
4. matrix clocks  
Order(N<sup>2</sup>) timestamps that say what A knows about what B knows about events on C.  
Acknowledgement vectors: an O(N) approximation to matrix clocks.

DUKE  
Systems & Architecture

## Update Ordering

**Problem:** how to ensure that all sites recognize a fixed order on updates, even if updates are delivered out of order?

**Solution:** Assign timestamps to updates at their accepting site, and order them by source timestamp at the receiver.

Assign nodes unique IDs: break ties with the origin node ID.

- What (if any) ordering exists between updates accepted by different sites?  
Comparing physical timestamps is arbitrary: physical clocks drift.  
Even a protocol to maintain loosely synchronized physical clocks cannot assign a meaningful ordering to events that occurred at “almost exactly the same time”.
- In Bayou, received updates may affect generation of future updates, since they are immediately visible to the user.

DUKE  
Systems & Architecture


## Causality and Logical Time

**Constraint:** The update ordering must respect *potential causality*.

- Communication patterns establish a **happened-before** order on events, which tells us when ordering *might* matter.
- Event  $e_1$  **happened-before**  $e_2$  iff  $e_1$  could possibly have affected the generation of  $e_2$ : we say that  $e_1 < e_2$ .  
 $e_1 < e_2$  iff  $e_1$  was “known” when  $e_2$  occurred.  
Events  $e_1$  and  $e_2$  are *potentially causally related*.
- In Bayou, users or applications may perceive inconsistencies if causal ordering of updates is not respected at all replicas.  
An update  $u$  should be ordered after all updates  $w$  known to the accepting site at the time  $u$  was accepted.  
e.g., the newsgroup example in the text.

DUKE  
Systems & Architecture

## Causality: Example



DUKE  
Systems & Architecture

## Logical Clocks

**Solution:** timestamp updates with *logical clocks* [Lamport]

Timestamping updates with the originating node's logical clock  $LC$  induces a partial order that respects potential causality.

**Clock condition:**  $e_1 < e_2$  implies that  $LC(e_1) < LC(e_2)$

1. Each site maintains a monotonically increasing clock value  $LC$ .
2. Globally visible events (e.g., updates) are timestamped with the current  $LC$  value at the generating site.


Increment local  $LC$  on each new event:  $LC = LC + 1$

3. Piggyback current clock value on all messages.

Receiver resets local  $LC$ : if  $LC_r > LC_s$  then  $LC_r = LC_s + 1$

DUKE  
Systems & Architecture

## Logical Clocks: Example



DUKE  
Systems & Architecture

## Flooding and the Prefix Property

In Bayou, each replica's knowledge of updates is determined by its pattern of communication with other nodes.

Loosely, a site knows everything that it *could* know from its contacts with other nodes.

- Anti-entropy *floods* updates.  
Tag each update originating from site  $i$  with *accept stamp* ( $i, LC_i$ ).  
Updates from each site are bulk-transmitted *cumulatively* in an order consistent with their source accept stamps.
- Flooding guarantees the *prefix property* of received updates.  
If a site knows an update  $u$  originating at site  $i$  with accept stamp  $LC_u$ , then it also knows all preceding updates  $w$  originating at site  $i$ : those with accept stamps  $LC_w < LC_u$ .

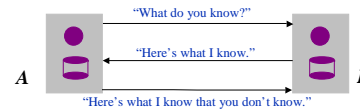
DUKE  
Systems & Architecture

## Which Updates to Propagate?

In an anti-entropy exchange,  $A$  must send  $B$  all updates known to  $A$  that are not yet known to  $B$ .

**Problem:** which updates are those?

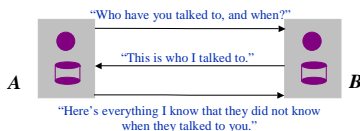
one-way "push" anti-entropy exchange (Bayou reconciliation)



DUKE  
Systems & Architecture

## Causality and Reconciliation


In general, a transfer from  $A$  must send  $B$  all updates that did not *happen-before* any update known to  $B$ .



Can we determine which updates to propagate by comparing logical clocks  $LC(A)$  and  $LC(B)$ ? **NO.**

DUKE  
Systems & Architecture

## Causality and Updates: Example



DUKE  
Systems & Architecture

### Motivation for Vector Clocks

Logical clocks induce an order consistent with causality, but it is actually stronger than causality.

- The converse of the *clock condition* does not hold: it may be that  $LC(e_1) < LC(e_2)$  even if  $e_1$  and  $e_2$  are concurrent.

If  $A$  could know anything  $B$  knows, then  $LC_A > LC_B$ , but if  $LC_A > LC_B$  then that doesn't make it so: "false positives".

Concurrent updates may be ordered unnecessarily.

- In Bayou, logical clocks tell us that  $A$  has *not* seen any update  $u$  with  $LC(u) > LC_A$ , but what has  $A$  seen?

$LC_A$  does not say if  $A$  saw a given update  $w$  with  $LC(w) < LC_A$ .

We need a clock mechanism that is not so sloppy about capturing causality.

DUKE  
Systems & Architecture

### Vector Clocks

Vector clocks (aka *vector timestamps* or *version vectors*) are a more detailed representation of what a site might know.

- In a system with  $N$  nodes, each site keeps a vector timestamp  $TS[N]$  as well as a logical clock  $LC$ .

$TS[j]$  at site  $i$  is the most recent value of site  $j$ 's logical clock that site  $i$  "heard about".

$TS[i] = LC_i$ ; each site  $i$  keeps its own  $LC$  in  $TS[i]$ .

- When site  $i$  generates a new event, it increments its logical clock.

$TS[i] = TS[i] + 1$

- A site  $r$  observing an event (e.g., receiving a message) from site  $s$  sets its  $TS_r$  to the pairwise maximum of  $TS_s$  and  $TS_r$ .

For each site  $i$ ,  $TS_r[i] = \text{MAX}(TS_s[i], TS_r[i])$

DUKE  
Systems & Architecture


### Vector Clocks and Causality

Vector clocks induce an order that *exactly* reflects causality.

- Tag each event  $e$  with current  $TS$  vector at originating site.  
vector timestamp  $TS(e)$
- $e_1$  **happened-before**  $e_2$  if and only if  $TS(e_2)$  **dominates**  $TS(e_1)$   
 $e_1 < e_2$  iff  $TS(e_1)[i] \leq TS(e_2)[i]$  for each site  $i$   
"Every event or update visible when  $e_1$  occurred was also visible when  $e_2$  occurred."
- Vector timestamps allow us to ask if two events are concurrent, or if one **happened-before** the other.  
If  $e_1 < e_2$  then  $LC(e_1) < LC(e_2)$  and  $TS(e_2)$  dominates  $TS(e_1)$ .  
If  $TS(e_2)$  does **not** dominate  $TS(e_1)$  then it is not true that  $e_1 < e_2$ .

DUKE  
Systems & Architecture

### Vector Clocks: Example



DUKE  
Systems & Architecture

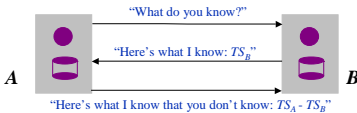
### Reconciliation with Vector Clocks

$A$  can determine which updates to pass to  $B$  by comparing their current vector clocks.

Tag each update originating from site  $i$  with *accept stamp*  $(i, LC_i)$ .

If  $TS_A[i] > TS_B[i]$ , then  $A$  has updates from  $i$  that  $B$  has not seen, i.e., they did not **happen-before** any updates known to  $B$ .

$A$  sends  $B$  all updates from site  $i$  tagged with accept stamps  $LC$  such that  $TS_B[i] < LC \leq TS_A[i]$ .



DUKE  
Systems & Architecture

### The Prefix Property and Reconciliation

Vector clocks work for anti-entropy transfers because they precisely encapsulate which updates a site has seen.

- The *prefix property* must hold for this to work.

If a site knows an update  $u$  originating at site  $i$  with accept stamp  $LC_u$ , then it also knows all preceding updates  $w$  originating at site  $i$ ; those with accept stamps  $LC_w < LC_u$ .

List-Group">

- $TS_B[i]$  exceeds the origin timestamp ( $LC$ ) of the *latest* update generated by  $i$  and received at  $B$ .

List-Group">

- Updates  $w$  from  $i$  with origin timestamps  $LC(w) > TS_B[i]$  are exactly those updates that did not **happen-before**  $TS_B$ .

If  $LC(w) > TS_B[i]$ , then  $TS_B$  cannot dominate  $TS_i(w)$ , so  $w$  cannot be known to  $B$ .

(Converse left as an exercise.)

DUKE  
Systems & Architecture

### When to Discard or Stabilize Updates?

**Problem 1:** when can a site discard its pending updates?

When can  $A$  know that every other site has seen some update  $u$ ?

**Problem 2:** when to commit (stabilize) pending updates?

A committed update  $w$  is *stable*: no other update  $u$  can arrive with  $u < w$ ;  $w$  will never need to be rolled back.

These two questions are equivalent:

Suppose we know that each peer site  $s$  has received this update  $w$ .

We had transitive contact with  $s$  since it received  $w$ .

We have received all updates known by  $s$  when it received  $w$ .

We have received all updates generated by  $s$  before it received  $w$ .

We have received all updates generated before  $w$ :  $w$  is stable.

One approach: propagate “knowledge about knowledge” with the AE.

DUKE  
Systems & Architecture

### The Need for Propagating Acknowledgments

Vector clocks tell us what  $B$  knows about  $C$ , but they do not reflect what  $A$  knows about what  $B$  knows about  $C$ .

Nodes need this information to determine when it is safe to discard/stabilize updates.

- $A$  can always tell if  $B$  has seen an update  $u$  by asking  $B$  for its vector clock and looking at it.

If  $u$  originated at site  $i$ , then  $B$  knows about  $u$  if and only if  $TS_B$  covers its accept stamp  $LC_i$ :  $TS_B[i] \geq LC_i$ .

- $A$  can only know that every site has seen  $u$  by looking at the vector clocks for every site.

Even if  $B$  recently received updates from  $C$ ,  $A$  cannot tell (from looking at  $B$ 's vector clock) if  $B$  got  $u$  from  $C$  or if  $B$  was already aware of  $u$  when  $C$  contacted it.

DUKE  
Systems & Architecture

### Solution 1: Matrix Clocks

Matrix clocks extend vector clocks to capture “what  $A$  knows about what  $B$  knows about  $C$ ”.

- Each site  $i$  maintains a matrix  $MC_i(N, N)$ .

Row  $j$  of  $i$ 's matrix clock  $MC_i$  is the most recent value of  $j$ 's vector clock  $TS_j$  that  $i$  has heard about.

$MC_i[i, i] = LC_i$  and  $MC_i[i, *] = TS_i$

$MC_i[j, k]$  = what  $i$  knows about what  $j$  knows about what happened at  $k$ .

- If  $A$  sends a message to  $B$ , then  $MC_B$  is set to the pairwise maximum of  $MC_A$  and  $MC_B$ .

If  $A$  knows that  $B$  knows  $u$ , then after  $A$  talks to  $C$ ,  $C$  knows that  $B$  knows  $u$  too.

DUKE  
Systems & Architecture

### Solution 2: Acknowledgment Stamps

Matrix clocks require  $N^2$  state for a system with  $N$  nodes.

Propagate  $N^2$  state on each exchange.

For anti-entropy, we can conservatively approximate the matrix clock using only  $N^2$  state. [Golding]

- After  $A$  passes updates to  $B$ , compute  $B$ 's *ack stamp* as the lowest  $LC$  entry in  $TS_{B-1}$ .

- Each node keeps an *acknowledgment summary vector*  $AS(N)$  of ack stamps for every other node “last it heard”.

( $AS$  vector just has the min value of each row in the matrix clock.)

- In an anti-entropy exchange from  $A$  to  $B$ , compute  $AS_B$  as a pairwise maximum of  $AS_A$  and  $AS_B$ .

$AS[j, j]$  does not tell  $i$  what  $j$  knows about  $k$ , but it does say that  $j$  knows about every event at  $k$  prior to  $AS[j, j]$ , for every  $k$ .

DUKE  
Systems & Architecture


### Golding's TSAE Protocol

Golding defined a Timestamped Anti-Entropy (TSAE) protocol that predates Bayou.

- designed for replicated Internet services (e.g., *refdbms*)
- reconciliation by two-way pairwise anti-entropy exchanges flooded updates
- studied role of network topology in partner selection
- uses logical clocks for accept stamps total commit ordering defined by logical clocks
- propagates knowledge of peer replica state by ack stamps and ack vectors

DUKE  
Systems & Architecture

### Ack Vectors in TSAE: Example



DUKE  
Systems & Architecture

## The Trouble With Ack Vectors

Matrix clocks and ack vectors can impede forward progress in the presence of failures.

- If a replica  $A$  fails or becomes disconnected, other nodes recognize that it is "getting behind".  
No node's ack stamp can advance beyond the accept stamp  $LC(w)$  of the last update  $w$  received by  $A$ .
- If a replica gets behind, other nodes cannot retire (discard) received updates  $u$  with  $LC(u) > LC(w)$ .
- One solution is to forcibly remove the disconnected node  $A$  from the replica set.

How to bring  $A$  up to date if it later rejoins the replica set? How to order updates generated by  $A$  while disconnected?

DUKE  
Systems & Architecture

## Committing Updates in Bayou

Bayou commits updates more aggressively using a *primary-commit protocol*.

- A single site is designated as the *primary*.
- The primary *commits* updates as it receives them.  
Primary assigns a *commit sequence number (CSN)* to each committed update.  
The final total update order is defined by CSN order.
- Sites learn of commitment through anti-entropy transfers.  
A site may learn of an update before learning that the update has been committed by the primary.  
Sites learn of commitment in CSN order. Updates known to be committed are *stable*: their order will never change.

DUKE  
Systems & Architecture

## Reconciliation with CSNs

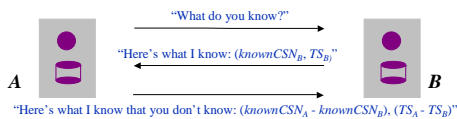
Each site also maintains a *knownCSN* counter, the CSN of the latest committed update the site knows has committed.

In an anti-entropy transfer,  $A$  looks at  $B$ 's *knownCSN*.

If  $A$  knows update  $w$  has committed, and  $CSN(w) > knownCSN_B$ , then  $A$  notifies  $B$  that  $w$  has committed.


$B$  updates *knownCSN<sub>B</sub>* as it learns of committed updates.

(This assumes that sites learn of commitment in CSN order.)



DUKE  
Systems & Architecture

## Bayou Update Logs: Example




ordering properties (so far)

- Each site sees/logs its own updates in accept order.
- Each site sees/logs updates from each peer site in accept order.
- Each site sees all updates in a causal order.
- Each site reorders received updates to total accept stamp order.
- Each site sends updates in its update log order.

DUKE  
Systems & Architecture

## Update Log with Commitment: Example



Suppose  $C$  is the primary

- $C$  commits updates in its update-log order.
- New constraint: all sites order known-committed updates before all tentative updates.
- Sites propagate updates/knowledge in update log order: known-committed updates propagate before tentative updates, and commit knowledge propagates in CSN order.
- Can CSN order violate causality? Can it violate a total order based on accept stamps?

DUKE  
Systems & Architecture

## Discarding Updates in Bayou

Any site  $A$  may truncate any prefix of the stable (committed) portion of its update log to reclaim space.

- $A$  needs no record of known-committed updates for itself.  
Committed updates are never rolled back, since commits are received in CSN order, the same as the final order.
- $A$  keeps stable updates in its log only so it can tell other sites about those updates.
- How can  $A$  reconcile with peers needing discarded updates?  
Easy: may discard stable updates if there is some other way to reconcile, e.g., send entire committed database state before sending updates.  
Truncation is a tradeoff: it reclaims local storage, but may make later reconciliations more expensive.

DUKE  
Systems & Architecture



### Detecting Update Conflicts: Traditional View

Many systems view updates  $u$  and  $w$  as *nonconflicting* iff:

- $u$  and  $w$  update different objects (e.g., files or records)
- $u$  and  $w$  update the same object, but one writer had observed the other's update before making its own.

In other words, an update conflict occurs when two processes  $p$  and  $q$  generate "concurrent" updates to the same object.

- $p$  and  $q$  updated the same object, but neither update *happened-before* the other.
- Updates to an object must follow a well-defined causal chain.  
Potential causality must induce a total order on the updates.

DUKE  
Systems & Architecture

### Example: Coda

Coda is a highly-available replicated file system (successor to AFS) that supports disconnected operation.

- Data is stored in *files*, which are grouped in *volumes*, which are stored on *servers*.

*Files are the granularity of caching by clients.*

- Volumes are the granularity of replication.

*VSG = Volume Storage Group == set of servers for a volume.*

- Availability by *read-any/write-all-available* replication.

*On an open, read the file from the "most up-to-date" member of the Available VSG (AVSG).*

*On a close after write, send the new version of the file to all members of the AVSG.*

DUKE  
Systems & Architecture

### Version Vectors in Coda

How to characterize the "up-to-dateness" of a file version?

Solution: Coda Version Vectors.

- Coda nodes maintain a version vector *CVV* for each file  $F$ .  
*CVV* has one element for each server in the file's VSG.  
*CVV[i]* is the # of writes received on *this version* of  $F$  at server  $i$ .
- On an *open*, client sets *CVV* to the server's *CVV*.
- On a *close*, client updates *CVV* and propagates it to the AVSG.  
*Increment CVV[i] for each server  $i$  that acknowledges the write.*
- We can compare the *CVVs* to tell if one version of  $F$  has updates not reflected in the other.  
*Two versions conflict if neither CVV dominates the other.*

DUKE  
Systems & Architecture

### Update Conflicts: the Bayou Approach

Bayou rejects the traditional "blind" view of conflicts:

- Updates might conflict even if they affect different records.  
*Example: two meeting-room records that contain the same room number and overlapping times.*  
*Example: two bibliography database entries that describe the same document.*  
*Example: two bibliography database entries that describe different documents using the same tag.*
- Concurrent updates to the same record might not conflict.  
*Writes don't conflict if they commute, e.g., they update different fields of the same record.*

Detecting/resolving conflicts is *application-specific*.

DUKE  
Systems & Architecture

### Application-Specific Reconciliation

Only the application can decide how to reconcile conflicting updates detected as writes are applied.

E.g., *refdbms*.

- Discard updates and deletions for already-deleted records.
- Change entry tag names to resolve add/add conflicts.  
*e.g., change Lampport75 to Lampport75a*
- field-specific update conflict resolution

DUKE  
Systems & Architecture

### Handling Update Conflicts in Bayou

- The primary commit protocol ensures that all sites commit the same writes in the same order.
- But this is not sufficient to guarantee that replicas converge.  
*Dependency checks and merge procedures handle conflicts.*  
*Check/merge can examine any state in the replica.*
- Check and merge procedures must be deterministic.  
*Limit inputs to the current contents of the database.*  
*Execute with fixed resource bounds so they fail deterministically.*
- Check/merge is executed every time a write is applied.  
*Rollback must be able to undo the effects of a merge.*

DUKE  
Systems & Architecture